

SUDHANSU KAKKAR

Phone: 732-500-2147 | Email: emsudhanshu@gmail.com | LinkedIn: linkedin.com/in/sudhansu-kakkar |
Portfolio: emsudhanshu.github.io/portfolio | GitHub: github.com/emsudhanshu

SUMMARY

Applied AI Engineer skilled in NLP, LLM fine-tuning, RAG, ML pipelines, with 5+ years of full-stack engineering experience.

TECHNICAL PROFICIENCY

- AI / ML:** Python, PyTorch, TensorFlow, scikit-learn, Transformers, HuggingFace, BERT Fine-Tuning, NLP, Text Classification, LLMs, RAG Pipelines, Vector Databases (FAISS), LlamaIndex, LangChain, CrewAI, Prompt Engineering, Model Evaluation (F1, ROC-AUC, BLEU, ROUGE), Streamlit, Time-Series Forecasting (Prophet), Anomaly Detection
- Backend & Systems:** Java Spring Boot, Node.js, Microservices, REST APIs, PostgreSQL, MongoDB, AWS, CI/CD, Git, Apache Spark, ETL Pipelines
- Frontend:** React, Redux Toolkit, HTML, CSS, JavaScript, Material UI, Micro-Frontends

RELEVANT EXPERIENCE

- Research Assistant - AI/ML** — Stevens Institute of Technology — May 2025 to Present
 - Designed and fine-tuned transformer (BERT) to classify multi-level depression from text, achieving 81% accuracy.
 - Built complete ML pipelines including dataset gathering, cleaning, preprocessing, tokenization, model training, validation, and hyperparameter tuning.
 - Implemented evaluation workflows (F1-score, ROC-AUC, confusion analysis) and performed targeted error analysis to improve model stability and prediction consistency.
- AI Engineer and Developer** — OncRef — Jun 2025 to Jul 2025
 - Built autonomous CrewAI agent workflows to automate web-scraping and summarization with minimal human supervision.
 - Converted unstructured research webpages into structured JSON through LLM-based enrichment and classification steps to improve data usability and accuracy.
 - Developed a lightweight end-to-end automation pipeline combining scraping, cleaning, and LLM-powered processing, reducing manual data preparation for the team.
- Senior SDE - Full Stack Developer** — Fidelity Information Services (FIS) — Nov 2021 to Jan 2025
 - Delivered major internet-banking features and standalone React + Spring Boot applications, supporting 400+ daily transactions and integrating with 15+ backend microservices across mission-critical systems.
 - Improved engineering efficiency by introducing reusable UI components, micro-frontend patterns, and Node.js automation scripts, reducing manual build effort by 2+ hours/day and accelerating UI development cycles by ~50%.
 - Enhanced system reliability and security by optimizing architecture, enforcing structured code reviews, and contributing to releases with reduced defect rates and strong 94% AppSec compliance.
- System Engineer - Front End Developer** — Tata Consultancy Services (TCS) — Aug 2019 to Nov 2021
 - Developed a cross-platform logistics app using React Native, Redux, and Material UI in collaboration with product teams.
 - Ensured performance and stability through testing, debugging, deployment, and defect resolution across platforms.
 - Improved user experience by fixing 200+ bugs and presenting weekly progress demos to stakeholders.

EDUCATION

- Master's in Applied Artificial Intelligence** — Stevens Institute of Technology, United States — 2025 to 2026 (Expected)
- Bachelor's in Computer Science Engineering (CGPA: 8.4)** — GGSIPU, India — 2015 to 2019

PROJECTS

- TruthfulQA Evaluation Pipeline — LLM Truthfulness Scoring** — (2025)
 - Evaluated LLM factual accuracy on TruthfulQA using GPT-4o-mini answers, BLEU/ROUGE/BLEURT metrics, and a DistilBERT truth classifier, achieving ~84% truthfulness.
- Multi-Class Depression Detection using BERT** — (2025)
 - Transformer-based multi-class classifier predicting depression severity from social media text with ~81% accuracy.
- Air Quality Index Forecasting & Anomaly Detection** — (2025)
 - Processed 29k+ AQI records with a Spark pipeline. Used Prophet to forecast PM2.5 with automated anomaly detection.
- AI Resume Chatbot Twin (RAG + Llama 3)** — (2025)
 - Llama-3 based RAG chatbot using all-MiniLM embeddings and JSON resume data stored in a persisted vector index.
- Dynamic Attendance Manager** — (2019)
 - Role-based attendance system using React, Redux-Saga, and Material UI for students, teachers, and admins.

CERTIFICATIONS AND TRAININGS

- Generative AI: Working with Large Language Models** — LinkedIn Learning (2025)
- AI Coding Agents with GitHub Copilot & Cursor** — LinkedIn Learning (2025)

ADDITIONAL SKILLS

Keywords: LLM Evaluation (BLEURT, Confusion Matrix), Error Analysis, Semantic Search Optimization, RAG Memory Optimization, Vector Search, Embedding Quality Tuning, Prompt Engineering, JSON Knowledge Structuring, Model Debugging, Data Cleaning & Text Normalization, Web-Scraping Automation, Production ML Deployment (FastAPI, Docker, CI/CD)